

sklearn.linear_model.LogisticRegression

```
class sklearn.linear_model.LogisticRegression(penalty='l2', *, dual=False, tol=0.0001, C=1.0, fit_intercept=True,
intercept_scaling=1, class_weight=None, random_state=None, solver='lbfgs', max_iter=100, multi_class='auto', verbose=0,
warm_start=False, n_jobs=None, l1_ratio=None)
```

[\[source\]](#)

Logistic Regression (aka logit, MaxEnt) classifier.

In the multiclass case, the training algorithm uses the one-vs-rest (OvR) scheme if the ‘multi_class’ option is set to ‘ovr’, and uses the cross-entropy loss if the ‘multi_class’ option is set to ‘multinomial’. (Currently the ‘multinomial’ option is supported only by the ‘lbfgs’, ‘sag’, ‘saga’ and ‘newton-cg’ solvers.)

This class implements regularized logistic regression using the ‘liblinear’ library, ‘newton-cg’, ‘sag’, ‘saga’ and ‘lbfgs’ solvers. **Note that regularization is applied by default.** It can handle both dense and sparse input. Use C-ordered arrays or CSR matrices containing 64-bit floats for optimal performance; any other input format will be converted (and copied).

The ‘newton-cg’, ‘sag’, and ‘lbfgs’ solvers support only L2 regularization with primal formulation, or no regularization. The ‘liblinear’ solver supports both L1 and L2 regularization, with a dual formulation only for the L2 penalty. The Elastic-Net regularization is only supported by the ‘saga’ solver.

Read more in the [User Guide](#).

Parameters:

penalty : {'l1', 'l2', 'elasticnet', None}, default='l2'

Specify the norm of the penalty:

- None: no penalty is added;
- 'l2': add a L2 penalty term and it is the default choice;
- 'l1': add a L1 penalty term;
- 'elasticnet': both L1 and L2 penalty terms are added.

Warning: Some penalties may not work with some solvers. See the parameter `solver` below, to know the compatibility between the penalty and solver.

New in version 0.19: l1 penalty with SAGA solver (allowing 'multinomial' + L1)

dual : bool, default=False

Dual (constrained) or primal (regularized, see also [this equation](#)) formulation. Dual formulation is only implemented for l2 penalty with liblinear solver. Prefer dual=False when n_samples > n_features.

tol : float, default=1e-4

Tolerance for stopping criteria.

C : float, default=1.0

Inverse of regularization strength; must be a positive float. Like in support vector machines, smaller values specify stronger regularization.

fit_intercept : bool, default=True

Specifies if a constant (a.k.a. bias or intercept) should be added to the decision function.

intercept_scaling : float, default=1

Useful only when the solver 'liblinear' is used and self.fit_intercept is set to True. In this case, x becomes [x, self.intercept_scaling], i.e. a “synthetic” feature with constant value equal to intercept_scaling is appended to the instance vector. The intercept becomes `intercept_scaling * synthetic_feature_weight`.

Note! the synthetic feature weight is subject to l1/l2 regularization as all other features. To lessen the effect of regularization on synthetic feature weight (and therefore on the intercept) intercept_scaling has to be increased.

class_weight : dict or 'balanced', default=None

Weights associated with classes in the form `{class_label: weight}`. If not given, all classes are supposed to have weight one.

The “balanced” mode uses the values of y to automatically adjust weights inversely proportional to class frequencies in the input data as `n_samples / (n_classes * np.bincount(y))`.

Note that these weights will be multiplied with sample_weight (passed through the fit method) if sample_weight is specified.

New in version 0.17: class_weight='balanced'

random_state : int, RandomState instance, default=None

Used when `solver` == 'sag', 'saga' or 'liblinear' to shuffle the data. See [Glossary](#) for details.

solver : {'lbfgs', 'liblinear', 'newton-cg', 'newton-cholesky', 'sag', 'saga'}, default='lbfgs'

Algorithm to use in the optimization problem. Default is 'lbfgs'. To choose a solver, you might want to consider the following aspects:

- For small datasets, 'liblinear' is a good choice, whereas 'sag' and 'saga' are faster for large ones;
- For multiclass problems, only 'newton-cg', 'sag', 'saga' and 'lbfgs' handle multinomial loss;
- 'liblinear' is limited to one-versus-rest schemes.
- 'newton-cholesky' is a good choice for `n_samples >> n_features`, especially with one-hot encoded categorical features with rare categories. Note that it is limited to binary classification and the one-versus-rest reduction for multiclass classification. Be aware that the memory usage of this solver has a quadratic dependency on `n_features` because it explicitly computes the Hessian matrix.

Warning: The choice of the algorithm depends on the penalty chosen. Supported penalties by solver:

- 'lbfgs' - ['l2', None]
- 'liblinear' - ['l1', 'l2']
- 'newton-cg' - ['l2', None]
- 'newton-cholesky' - ['l2', None]
- 'sag' - ['l2', None]
- 'saga' - ['l2', None]
- 'elasticnet' - ['l1', 'l2', None]

Toggle Menu

Note: ‘sag’ and ‘saga’ fast convergence is only guaranteed on features with approximately the same scale. You can preprocess the data with a scaler from [sklearn.preprocessing](#).

See also: Refer to the User Guide for more information regarding [LogisticRegression](#) and more specifically the [Table](#) summarizing solver/penalty supports.

New in version 0.17: Stochastic Average Gradient descent solver.

New in version 0.19: SAGA solver.

Changed in version 0.22: The default solver changed from ‘liblinear’ to ‘lbfgs’ in 0.22.

New in version 1.2: newton-cholesky solver.

max_iter : *int, default=100*

Maximum number of iterations taken for the solvers to converge.

multi_class : {‘auto’, ‘ovr’, ‘multinomial’}, *default=‘auto’*

If the option chosen is ‘ovr’, then a binary problem is fit for each label. For ‘multinomial’ the loss minimised is the multinomial loss fit across the entire probability distribution, *even when the data is binary*. ‘multinomial’ is unavailable when solver=‘liblinear’. ‘auto’ selects ‘ovr’ if the data is binary, or if solver=‘liblinear’, and otherwise selects ‘multinomial’.

New in version 0.18: Stochastic Average Gradient descent solver for ‘multinomial’ case.

Changed in version 0.22: Default changed from ‘ovr’ to ‘auto’ in 0.22.

verbose : *int, default=0*

For the liblinear and lbfgs solvers set verbose to any positive number for verbosity.

warm_start : *bool, default=False*

When set to True, reuse the solution of the previous call to fit as initialization, otherwise, just erase the previous solution. Useless for liblinear solver. See [the Glossary](#).

New in version 0.17: warm_start to support lbfgs, newton-cg, sag, saga solvers.

n_jobs : *int, default=None*

Number of CPU cores used when parallelizing over classes if multi_class=‘ovr’”. This parameter is ignored when the solver is set to ‘liblinear’ regardless of whether ‘multi_class’ is specified or not. None means 1 unless in a [joblib.parallel_backend](#) context. -1 means using all processors. See [Glossary](#) for more details.

l1_ratio : *float, default=None*

The Elastic-Net mixing parameter, with $0 \leq \text{l1_ratio} \leq 1$. Only used if penalty=‘elasticnet’. Setting l1_ratio=0 is equivalent to using penalty=‘l2’, while setting l1_ratio=1 is equivalent to using penalty=‘l1’. For $0 < \text{l1_ratio} < 1$, the penalty is a combination of L1 and L2.

Attributes:

classes_ : *ndarray of shape (n_classes,)*

A list of class labels known to the classifier.

coef_ : *ndarray of shape (1, n_features) or (n_classes, n_features)*

Coefficient of the features in the decision function.

coef_ is of shape (1, n_features) when the given problem is binary. In particular, when multi_class=‘multinomial’, coef_ corresponds to outcome 1 (True) and -coef_ corresponds to outcome 0 (False).

intercept_ : *ndarray of shape (1,) or (n_classes,)*

Intercept (a.k.a. bias) added to the decision function.

If fit_intercept is set to False, the intercept is set to zero. intercept_ is of shape (1,) when the given problem is binary. In particular, when multi_class=‘multinomial’, intercept_ corresponds to outcome 1 (True) and -intercept_ corresponds to outcome 0 (False).

n_features_in_ : *int*

Number of features seen during fit.

Toggle Menu

New in version 0.24.

feature_names_in_ : ndarray of shape (n_features_in_,)

Names of features seen during [fit](#). Defined only when `x` has feature names that are all strings.

New in version 1.0.

n_iter_ : ndarray of shape (n_classes,) or (1,)

Actual number of iterations for all classes. If binary or multinomial, it returns only 1 element. For liblinear solver, only the maximum number of iteration across all classes is given.

Changed in version 0.20: In SciPy <= 1.0.0 the number of lbfgs iterations may exceed `max_iter`. `n_iter_` will now report at most `max_iter`.

See also:

[SGDClassifier](#)

Incrementally trained logistic regression (when given the parameter `loss="log_loss"`).

[LogisticRegressionCV](#)

Logistic regression with built-in cross validation.

Notes

The underlying C implementation uses a random number generator to select features when fitting the model. It is thus not uncommon, to have slightly different results for the same input data. If that happens, try with a smaller `tol` parameter.

Predict output may not match that of standalone liblinear in certain cases. See [differences from liblinear](#) in the narrative documentation.

References

L-BFGS-B – Software for Large-scale Bound-constrained Optimization

Ciyou Zhu, Richard Byrd, Jorge Nocedal and Jose Luis Morales. <http://users.iems.northwestern.edu/~nocedal/lbfgsb.html>

LIBLINEAR – A Library for Large Linear Classification

<https://www.csie.ntu.edu.tw/~cjlin/liblinear/>

SAG – Mark Schmidt, Nicolas Le Roux, and Francis Bach

Minimizing Finite Sums with the Stochastic Average Gradient <https://hal.inria.fr/hal-00860051/document>

SAGA – Defazio, A., Bach F. & Lacoste-Julien S. (2014).

[“SAGA: A Fast Incremental Gradient Method With Support for Non-Strongly Convex Composite Objectives”](#)

Hsiang-Fu Yu, Fang-Lan Huang, Chih-Jen Lin (2011). Dual coordinate descent

methods for logistic regression and maximum entropy models. Machine Learning 85(1-2):41-75. https://www.csie.ntu.edu.tw/~cjlin/papers/maxent_dual.pdf

Examples

```
>>> from sklearn.datasets import load_iris
>>> from sklearn.linear_model import LogisticRegression
>>> X, y = load_iris(return_X_y=True)
>>> clf = LogisticRegression(random_state=0).fit(X, y)
>>> clf.predict(X[:2, :])
array([0, 0])
>>> clf.predict_proba(X[:2, :])
array([[9.8...e-01, 1.8...e-02, 1.4...e-08],
       [9.7...e-01, 2.8...e-02, ...e-08]])
>>> clf.score(X, y)
0.97...
```

Methods

| | |
|--|---|
| decision_function(X) | Predict confidence scores for samples. |
| densify() | Convert coefficient matrix to dense array format. |
| fit(X, y[, sample_weight]) | Fit the model according to the given training data. |
| get_metadata_routing() | Get metadata routing of this object. |
| get_params([deep]) | Get parameters for this estimator. |
| predict(X) | Predict class labels for samples in X. |
| _proba(X) | Predict logarithm of probability estimates. |

Toggle Menu

| | |
|--|---|
| <code>predict_proba(X)</code> | Probability estimates. |
| <code>score(X, y[, sample_weight])</code> | Return the mean accuracy on the given test data and labels. |
| <code>set_fit_request(*[, sample_weight])</code> | Request metadata passed to the <code>fit</code> method. |
| <code>set_params(**params)</code> | Set the parameters of this estimator. |
| <code>set_score_request(*[, sample_weight])</code> | Request metadata passed to the <code>score</code> method. |
| <code>sparsify()</code> | Convert coefficient matrix to sparse format. |

decision_function(X)

[source]

Predict confidence scores for samples.

The confidence score for a sample is proportional to the signed distance of that sample to the hyperplane.

Parameters:

X : {array-like, sparse matrix} of shape (n_samples, n_features)
The data matrix for which we want to get the confidence scores.

Returns:

scores : ndarray of shape (n_samples,) or (n_samples, n_classes)
Confidence scores per (n_samples, n_classes) combination. In the binary case, confidence score for `self.classes_[1]` where `>0` means this class would be predicted.

densify()

[source]

Convert coefficient matrix to dense array format.

Converts the `coef_` member (back) to a `numpy.ndarray`. This is the default format of `coef_` and is required for fitting, so calling this method is only required on models that have previously been sparsified; otherwise, it is a no-op.

Returns:

self
Fitted estimator.

fit(X, y, sample_weight=None)

[source]

Fit the model according to the given training data.

Parameters:

X : {array-like, sparse matrix} of shape (n_samples, n_features)
Training vector, where `n_samples` is the number of samples and `n_features` is the number of features.

y : array-like of shape (n_samples,)
Target vector relative to `X`.

sample_weight : array-like of shape (n_samples,) default=None
Array of weights that are assigned to individual samples. If not provided, then each sample is given unit weight.

New in version 0.17: sample_weight support to LogisticRegression.

Returns:

self
Fitted estimator.

Notes

The SAGA solver supports both float64 and float32 bit arrays.

get_metadata_routing()

[source]

Get metadata routing of this object.

Please check [User Guide](#) on how the routing mechanism works.

Returns:

routing : *MetadataRequest*
A [MetadataRequest](#) encapsulating routing information.

get_params(deep=True)

[source]

Get parameters for this estimator.

Parameters:

deep : *bool, default=True*
If True, will return the parameters for this estimator and contained subobjects that are estimators.

Returns:

params : *dict*
Parameter names mapped to their values.

predict(X)

[source]

Predict class labels for samples in X.

Parameters:

X : *{array-like, sparse matrix} of shape (n_samples, n_features)*
The data matrix for which we want to get the predictions.

Returns:

y_pred : *ndarray of shape (n_samples,)*
Vector containing the class labels for each sample.

predict_log_proba(X)

[source]

Predict logarithm of probability estimates.

The returned estimates for all classes are ordered by the label of classes.

Parameters:

X : *array-like of shape (n_samples, n_features)*
Vector to be scored, where `n_samples` is the number of samples and `n_features` is the number of features.

Returns:

T : *array-like of shape (n_samples, n_classes)*
Returns the log-probability of the sample for each class in the model, where classes are ordered as they are in `self.classes_`.

predict_proba(X)

[source]

Probability estimates.

The returned estimates for all classes are ordered by the label of classes.

For a multi_class problem, if multi_class is set to be “multinomial” the softmax function is used to find the predicted probability of each class. Else use a one-vs-rest approach, i.e. calculate the probability of each class assuming it to be positive using the logistic function. and normalize these values across all the classes.

Parameters:

X : *array-like of shape (n_samples, n_features)*
Vector to be scored, where `n_samples` is the number of samples and `n_features` is the number of features.

Returns:

T : *array-like of shape (n_samples, n_classes)*
Returns the probability of the sample for each class in the model, where classes are ordered as they are in `self.classes_`.

score(X, y, sample_weight=None)

[source]

Return the mean accuracy on the given test data and labels.

In multi-label classification, this is the subset accuracy which is a harsh metric since you require for each sample that each label set be correctly predicted.

Parameters:

- X : array-like of shape (n_samples, n_features)
- Test samples.
- y : array-like of shape (n_samples,) or (n_samples, n_outputs)
- True labels for `x`.
- sample_weight : array-like of shape (n_samples,), default=None
- Sample weights.

Returns:

- score : float
- Mean accuracy of `self.predict(X)` w.r.t. `y`.

set_fit_request(*, sample_weight: bool | None | str = '\$UNCHANGED\$') → LogisticRegression

[source]

Request metadata passed to the `fit` method.

Note that this method is only relevant if `enable_metadata_routing=True` (see [sklearn.set_config](#)). Please see [User Guide](#) on how the routing mechanism works.

The options for each parameter are:

- `True`: metadata is requested, and passed to `fit` if provided. The request is ignored if metadata is not provided.
- `False`: metadata is not requested and the meta-estimator will not pass it to `fit`.
- `None`: metadata is not requested, and the meta-estimator will raise an error if the user provides it.
- `str`: metadata should be passed to the meta-estimator with this given alias instead of the original name.

The default (`sklearn.utils.metadata_routing.UNCHANGED`) retains the existing request. This allows you to change the request for some parameters and not others.

New in version 1.3.

Note: This method is only relevant if this estimator is used as a sub-estimator of a meta-estimator, e.g. used inside a [Pipeline](#). Otherwise it has no effect.

Parameters:

- sample_weight : str, True, False, or None, default=sklearn.utils.metadata_routing.UNCHANGED
- Metadata routing for `sample_weight` parameter in `fit`.

Returns:

- self : object
- The updated object.

set_params(**params)

[source]

Set the parameters of this estimator.

The method works on simple estimators as well as on nested objects (such as [Pipeline](#)). The latter have parameters of the form `<component>__<parameter>` so that it's possible to update each component of a nested object.

Parameters:

****params : dict**
Estimator parameters.

Returns:

self : *estimator instance*
Estimator instance.

set_score_request(*, sample_weight: [bool](#) | [None](#) | [str](#) = '\$UNCHANGED\$') → [LogisticRegression](#) [\[source\]](#)

Request metadata passed to the `score` method.

Note that this method is only relevant if `enable_metadata_routing=True` (see [sklearn.set_config](#)). Please see [User Guide](#) on how the routing mechanism works.

The options for each parameter are:

- `True`: metadata is requested, and passed to `score` if provided. The request is ignored if metadata is not provided.
- `False`: metadata is not requested and the meta-estimator will not pass it to `score`.
- `None`: metadata is not requested, and the meta-estimator will raise an error if the user provides it.
- `str`: metadata should be passed to the meta-estimator with this given alias instead of the original name.

The default (`sklearn.utils.metadata_routing.UNCHANGED`) retains the existing request. This allows you to change the request for some parameters and not others.

New in version 1.3.

Note: This method is only relevant if this estimator is used as a sub-estimator of a meta-estimator, e.g. used inside a [Pipeline](#). Otherwise it has no effect.

Parameters:

sample_weight : str, True, False, or None, default=sklearn.utils.metadata_routing.UNCHANGED
Metadata routing for `sample_weight` parameter in `score`.

Returns:

self : object
The updated object.

sparsify() [\[source\]](#)

Convert coefficient matrix to sparse format.

Converts the `coef_` member to a `scipy.sparse` matrix, which for L1-regularized models can be much more memory- and storage-efficient than the usual `numpy.ndarray` representation.

The `intercept_` member is not converted.

Returns:

self
Fitted estimator.

Notes

For non-sparse models, i.e. when there are not many zeros in `coef_`, this may actually *increase* memory usage, so use this method with care. A rule of thumb is that the number of zero elements, which can be computed with `(coef_ == 0).sum()`, must be more than 50% for this to provide significant benefits.

After calling this method, further fitting with the `partial_fit` method (if any) will not work until you call `densify`.

Examples using `sklearn.linear_model.LogisticRegression`

Release Highlights for scikit-learn 1.3

Release Highlights for scikit-learn 1.1

Release Highlights for scikit-learn 1.0

Release Highlights for scikit-learn 0.24

Release Highlights for scikit-learn 0.23

Release Highlights for scikit-learn 0.22

Probability Calibration curves

Plot classification
probability

Feature transformation with ensembles of trees

Plot class probabilities
calculated by the
VotingClassifier

Model-based and sequential feature selection

Recursive feature elimination with cross-validation

Comparing various nonlinear solvers

L1 Penalty and Sparsity in Logistic Regression

Logistic Regression 3-class Classifier

Logistic function

MNIST classification
using multinomial logistic
+ L1

Multiclass sparse logistic
regression on
20newgroups

Plot multinomial and One-vs-Rest Logistic Regression

Regularization path of L1-Logistic Regression


Displaying Pipelines

Displaying estimators and complex pipelines


Introducing the set_output API

Visualizations with Display Objects


Class Likelihood Ratios
to measure classification
performance




Multiclass Receiver Operating Characteristic (ROC)




Multilabel classification using a classifier chain




Restricted Boltzmann Machine features for digit classification



Column Transformer with Mixed Types




Pipelining: chaining a PCA and a logistic regression



Feature discretization



Digits Classification Exercise



Classification of text documents using sparse features